

Olympus Analytics Project Overview

Olympus Analytics helps small and mid-sized businesses bring all their siloed data into one place. This project establishes a working data pipeline in Snowflake using simulated CRM and LMS data. The goal is to simulate our final product so we can show dashboards and AI chat insights to potential clients.

What You'll Build

Key Goals

- Set up a working Snowflake pipeline using dummy HubSpot and Disco LMS data
- Create dbt models for Bronze → Silver → Gold data layers
- Build views that power dashboards and chat-based analysis
- Build Dashboard + Chat feature in 1 end experience. Use tools like Retool or similar.
- Deliver clear, reusable code that we can connect to real systems later

Your Deliverables

- 1. Dummy Data Setup**
 - Create or generate CSVs that mimic HubSpot and Disco LMS APIs
 - Example HubSpot fields: contacts, deals, companies, activities
 - Example Disco fields: users, courses, completions, enrollments
 - Include basic fields like email, timestamps, deal stage, training status, and contract value
- 2. Bronze Layer (Raw)**
 - Load CSVs into Snowflake (as-is)
 - Store in `hubspot_raw` and `disco_raw` schemas
 - Add ingestion metadata
- 3. Silver Layer (Clean)**
 - Standardize field types (e.g. lowercase emails, UTC timestamps)
 - Remove duplicates and nulls
 - Create join keys (e.g. `person_id` from email)
 - Store in `hubspot_clean`, `disco_clean`, `shared_clean`
- 4. Gold Layer (Business)**
 - Design a clear star schema
 - Facts: deals, course completions, person-course links
 - Dimensions: person, course, account, time
 - Simulate realistic joins (e.g. people who took training → closed more deals)
- 5. Serving Layer**

- Create clean, readable views for BI and AI tools
 - Examples: `v_kpi_summary`, `v_rep_productivity`, `v_course_performance`
 - Views must be optimized and documented
6. **Testing & Validation**
- Use dbt to test for missing, duplicate, or invalid values
 - Simulate a second file drop to test incremental refresh
7. **Handoff Docs**
- Entity Relationship Diagram (ERD)
 - Field-level data dictionary
 - Step-by-step refresh/runbook
 - Loom walkthrough video
-

Tech Stack

- **Warehouse:** Snowflake
 - **Modeling:** dbt (with tests and docs)
 - **Orchestration:** dbt Cloud or Snowflake Tasks
 - **BI:** Hex (we'll plug it in)
 - **AI:** Snowflake Cortex (natural language → SQL)
 - **Source Data:** Simulated CSVs (you build)
-

Architecture Overview

Star Schema in Gold Layer

In the Gold Layer, use a star schema — one central fact table (like `fct_deal`) connects to descriptive dimension tables (like `dim_person`, `dim_time`).

- Facts = events and metrics
- Dimensions = who, what, when
- Keep joins simple and fast

Data Flow

1. CSVs → Snowflake Raw (Bronze)
2. Cleaned tables → Silver
3. Joined KPIs → Gold
4. Views → Cortex + BI

5. Manual refresh via dbt CLI or dbt Cloud
-



KPIs to Simulate

Organize KPIs into these business outcome categories:



Money Increased

- Contract value by course completion
- Deal win rate after training
- Revenue lift for trained reps
- Upsell success post-training



Time Saved

- Time to close first deal vs training date
- Sales cycle time by training progress
- Time to first qualified lead



Money Saved

- Activities per closed deal
- Cost per deal by rep training level
- (Optional) Tickets reduced per training

These should be visible in BI and AI chat outputs.



Ideal Qualifications

- 2+ years Snowflake + dbt
 - Built CRM or LMS data models before
 - Know star schema, SCDs, dimensional modeling
 - Clear communicator, documents well
-



Bonus Skills

- Familiar with Snowflake Cortex / GPT-based tools

- BI tool setup experience (Hex, Tableau, etc.)
 - CI/CD or dbt Cloud automation
 - Reverse ETL knowledge (e.g. Hightouch)
-

What We'll Provide

- Snowflake account (with access)
 - Airbyte workspace if needed
 - KPI guidance and check-ins
 - Slack or async comms
-

Timeline

- Start: ~Week of [Insert Date]
 - Duration: 2 weeks
 - Check-ins: Async + 1 call/week (30 mins) Available all day for questions
-

What Success Looks Like

- We can demo working dashboards powered by Snowflake
- Cortex can answer natural language questions from Gold Layer
- Code and docs are clean and ready for real data
- Olympus can show this to clients, investors, and engineers

Bonus Opportunity

If this MVP helps Olympus land its first paying client, you will become our engineer of choice moving forward. This means:

- Priority for future paid development work
- First in line as we scale into full production pipelines
- Ongoing opportunities to support new client builds and platform features

We're looking for someone to grow with us — deliver strong now and you'll be our go-to moving forward.

Let us know if anything is unclear — this SOW is your blueprint to help us move fast and show value.